

Arithmetic Algorithms and ALU Design

Mainak Chaudhuri

Indian Institute of Technology Kanpur

Sketch

- Abstract model of computer
- ALU architecture
- Arithmetic algorithms
 - Overflow detection in addition and subtraction
 - Integer multiplication
 - Integer division
 - Floating-point addition
 - Floating-point multiplication

Abstract model of computer


- Computer has an ISA
- The implementation of the ISA is an abstract five-state synchronous FSM
 - Each state change happens on posedge clock
 - State 0: fetch the instruction pointed to by program counter from memory; update program counter to point to the next instruction
 - State 1: decode the instruction to extract various fields and read source register operands
 - State 2: execute the instruction in ALU; compute address of load/store instructions; update program counter if control transfer instruction
 - State 3: access memory if load/store instruction
 - State 4: write result to destination register if the instruction produces a result

Abstract model of computer

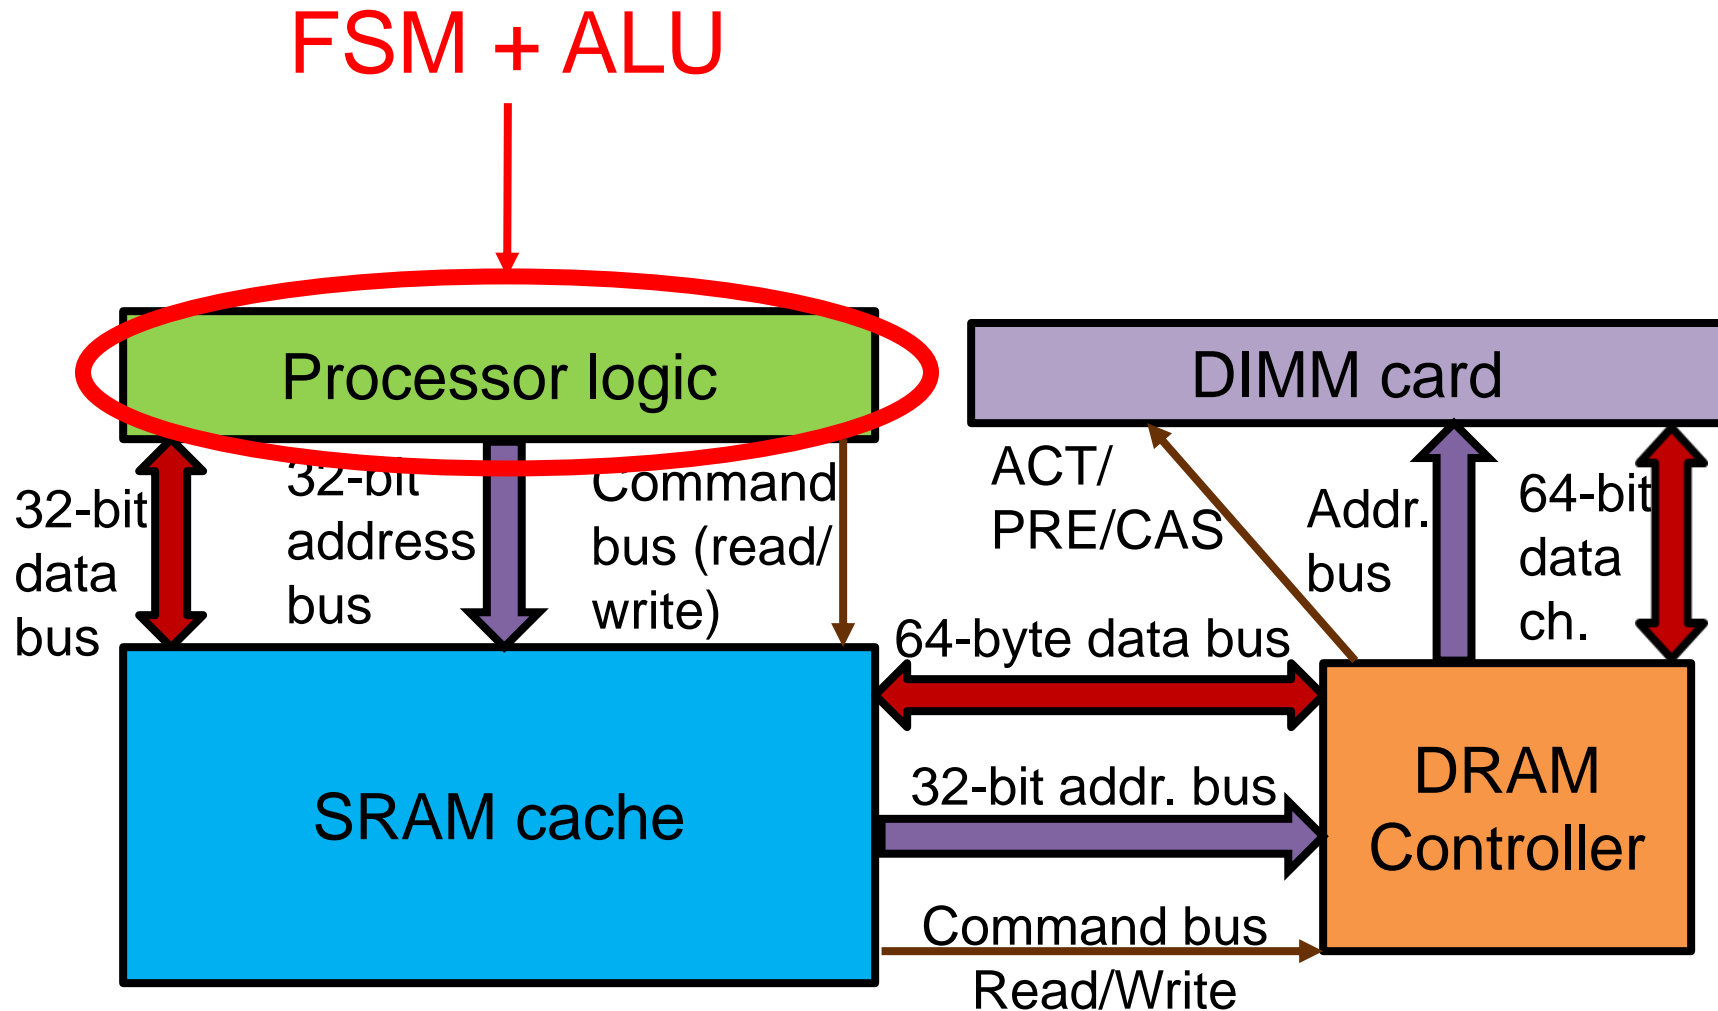
- Fetching an instruction requires **accessing memory** with the program counter as addr.
- Decoding an instruction for MIPS is simple due to small number of formats and fixed position of the register specifiers
- Reading operands from register file requires exercising the read ports
- Executing an instruction and computing address of a load/store instruction requires an **arithmetic logic unit (ALU)**
- Load/store instructions **access memory** with the computed address
- Writing result to register file requires exercising the write ports

Abstract model of computer

- The logic on processor chip implements the FSM and the ALU



Abstract model of computer



ALU architecture

- ALU is responsible for executing the core of the MIPS instructions
 - Everything except load/store instructions
- ALU takes two inputs for most instructions and produces a result that may or may not get written to a destination register
 - For example, control transfer instructions do not write to a general-purpose register, but writes to the program counter
- ALUs are of two kinds: integer and floating-point
 - The floating-point ALU is often referred to as the floating-point unit (FPU)

Integer ALU architecture

- Components of the integer ALU
 - Adder/subtractor
 - Executes add, addu, sub, subu, addi, addiu
 - Computes target of conditional branch instructions
 - Computes address of load/store instructions
 - Multiplier
 - Executes mult, multu
 - Divider
 - Executes div, divu
 - Logic gate array
 - Executes and, nor, or, xor, andi, ori, xori, lui
 - Comparator and branch subunit
 - Executes slt, sltu, slti, sltiu
 - Executes the comparison part of beq, bne, bgez, bgtz, bltz, blez

Integer ALU architecture

- Components of the integer ALU
 - Shifter
 - Executes sll, sllv, sra, srav, srl, srlv
 - Naïve implementation would shift by one bit position in every clock cycle
 - Inefficient for arbitrary shift amounts
 - Barrel shifters can shift an operand by arbitrary amount in time no more than an addition takes
 - Multiplexer (part of branch subunit)
 - Selects between branch target and PC+4 depending on the output of the comparison part of cond. branch
 - Data movement unit (part of multiplier/divider)
 - Executes mflo, mfhi, mtlo, mthi
 - Some designs do not consider the shifter, multiplier, and divider as part of the ALU

Overflow detection in MIPS

- MIPS ISA offers two types of arithmetic instructions for addition and subtraction
 - One type detects overflows (add, sub, addi) and the other doesn't (addu, subu, addiu)
 - Whether **overflow** should be detected in arithmetic operations is **a part of the HLL specification**
 - MIPS ISA can support both types of HLLs
 - If add/sub overflow is supposed to be detected, MIPS raises an exception on arithmetic overflow
 - Saves PC in EPC
 - **Saves all general-purpose and floating-point registers in memory**
 - Jumps to a fixed handler which examines cause and status registers to invoke the appropriate exception handler

Overflow detection in MIPS

- Arithmetic overflow exception
 - On completion of the handler
 - Restore all general-purpose and floating-point registers from memory
 - Compute EPC+4 and move it to \$k0 or \$k1 (same as \$26 and \$27)
 - These two registers are reserved for use of operating system (e.g., interrupt, exception, system call handlers)
 - These two registers are not preserved across handlers
 - Execute jr \$k0 or jr \$k1 depending on where EPC+4 is
- Overflow detection is also possible in software
 - Not very relevant for MIPS addition/subtraction
 - Important for processors that cannot detect overflow in hardware

Overflow detection in software

- Software cannot access the carry bits into and out of the most significant bit position
 - So, cannot compare these to detect overflow
- Observation
 - In signed addition, an overflow is detected if and only if the operands have the same sign and the result has sign opposite of the operand
 - If the operands have opposite signs, the result is guaranteed to have smaller magnitude than the larger magnitude operand and hence, there cannot be an overflow
 - In signed subtraction, an overflow is detected if and only if the operands have opposite signs and the result has sign opposite of the first operand
 - If the operands have the same sign, the result is guaranteed to have smaller magnitude than the larger magnitude operand

Overflow detection in software

- Code for detecting overflow in signed addition ($\$t0 = \$t1 + \$t2$)

```
addu $t0, $t1, $t2  # no overflow exception
```

```
xor  $t3, $t1, $t2
```

```
slt  $t3, $t3, $0    # $t3 is 1 iff signs of $t1, $t2 differ
```

```
bne  $t3, $0, no_overflow
```

```
xor  $t3, $t0, $t1
```

```
slt  $t3, $t3, $0    # $t3 is 1 iff signs of $t0, $t1 differ
```

```
bne  $t3, $0, overflow_handler
```

no_overflow: ...

Overflow detection in software

- Unsigned addition and subtraction
 - Overflow is detected if and only if the magnitude of the result exceeds the maximum representable unsigned value
 - $2^{32} - 1$ for 32-bit processors
- Code for detecting overflow in unsigned addition ($\$t0 = \$t1 + \$t2$)

```
addu $t0, $t1, $t2  # no overflow exception
```

```
nor  $t3, $t1, $0    # $t3 has  $2^{32} - \$t1 - 1$ 
```

```
sltu $t3, $t3, $t2   # $t3 has 1 iff  $\$t1 + \$t2 > 2^{32} - 1$ 
```

```
bne  $t3, $0, overflow_handler
```

- Note that the comparison must be unsigned (sltu) so that $2^{32} - \$t1 - 1$ and $\$t2$ are treated as unsigned operands

Multiplication algorithm

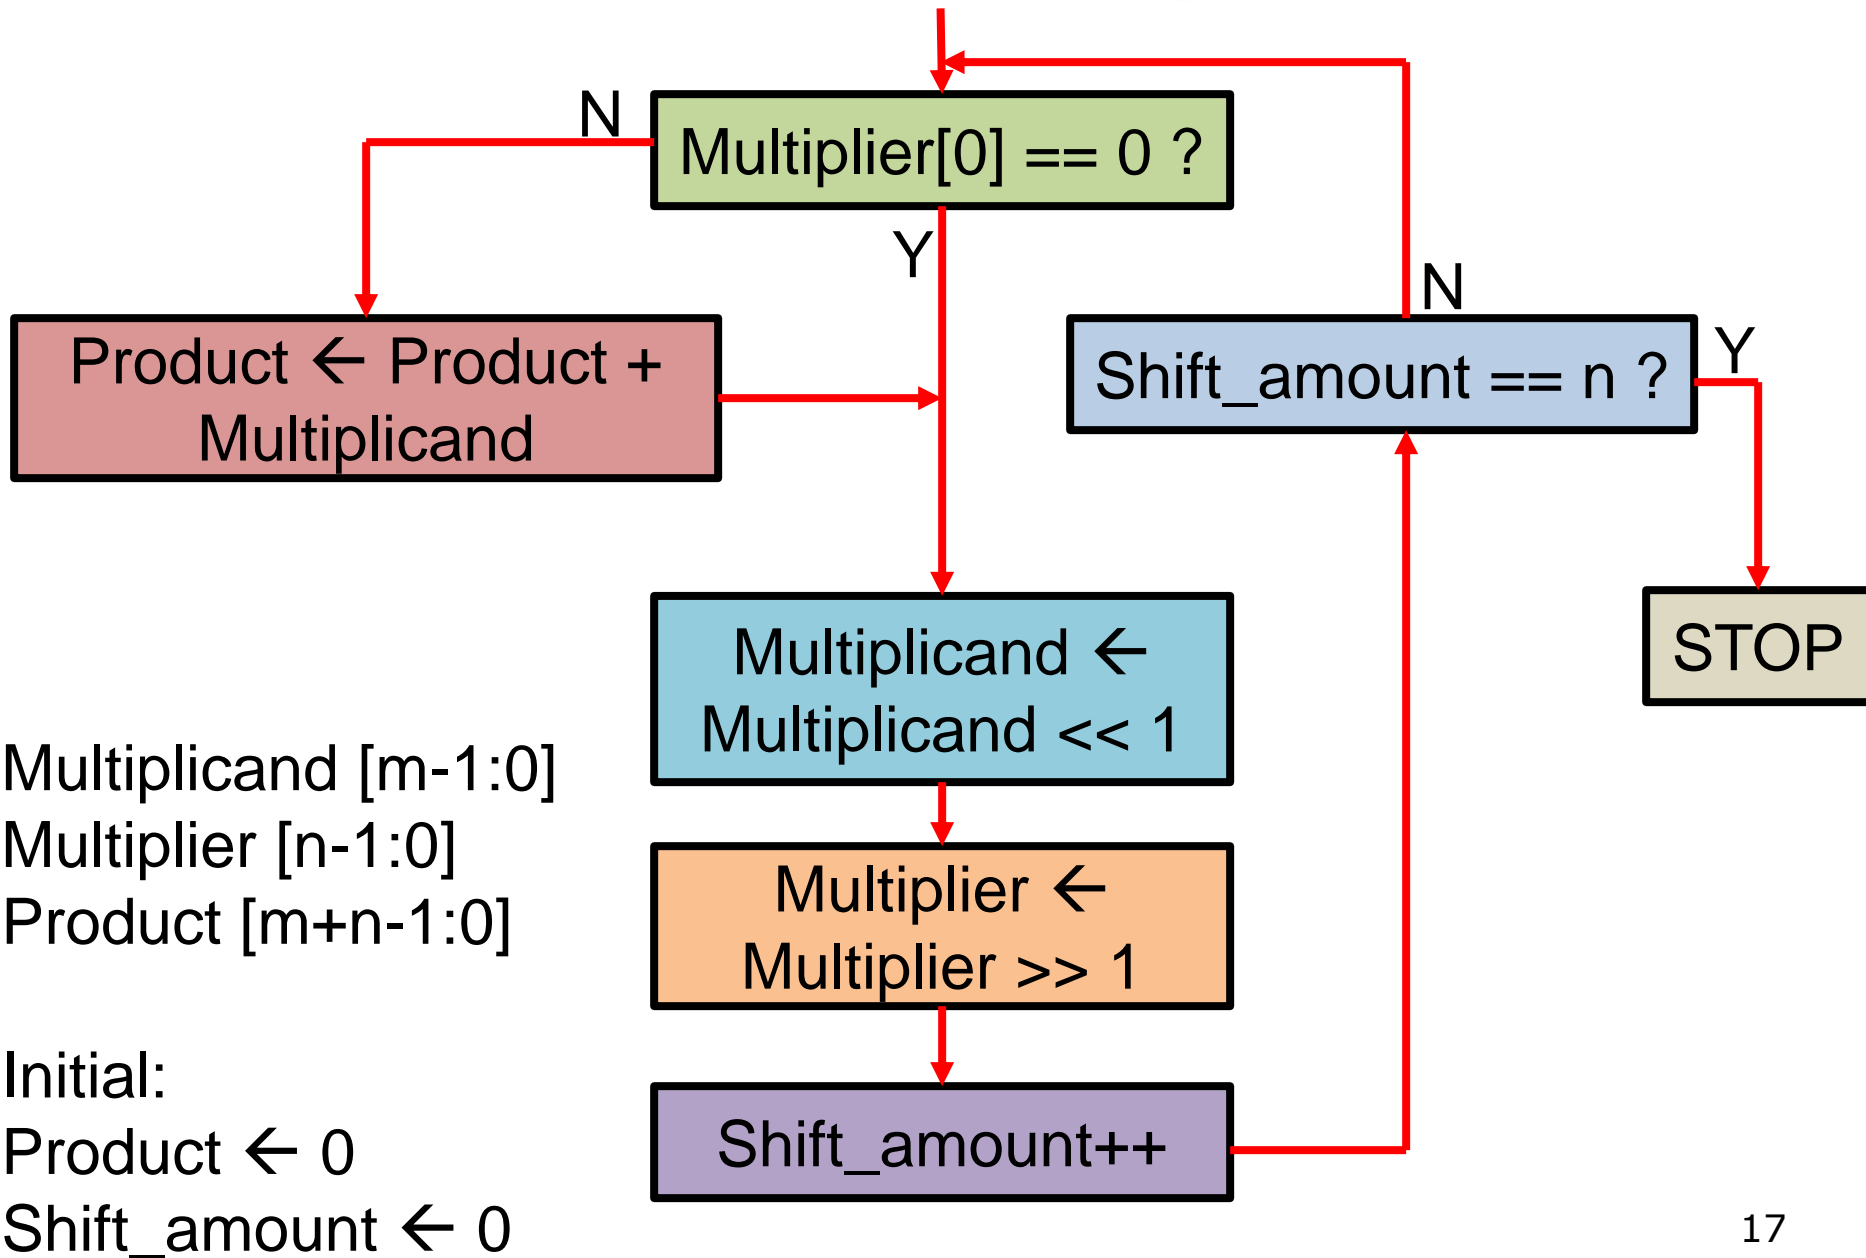
- One operand is called the multiplicand and the other is called the multiplier
 - m-bit multiplicand and n-bit multiplier lead to (m+n)-bit product
- Paper-pencil algorithm (assume unsigned operands)

$$\begin{array}{r} 1001 \text{ multiplicand} \\ \times 1010 \text{ multiplier} \\ \hline 0000 \\ 1001 \text{ (why shift left?)} \\ 0000 \\ 1001 \\ \hline 1011010 \end{array}$$

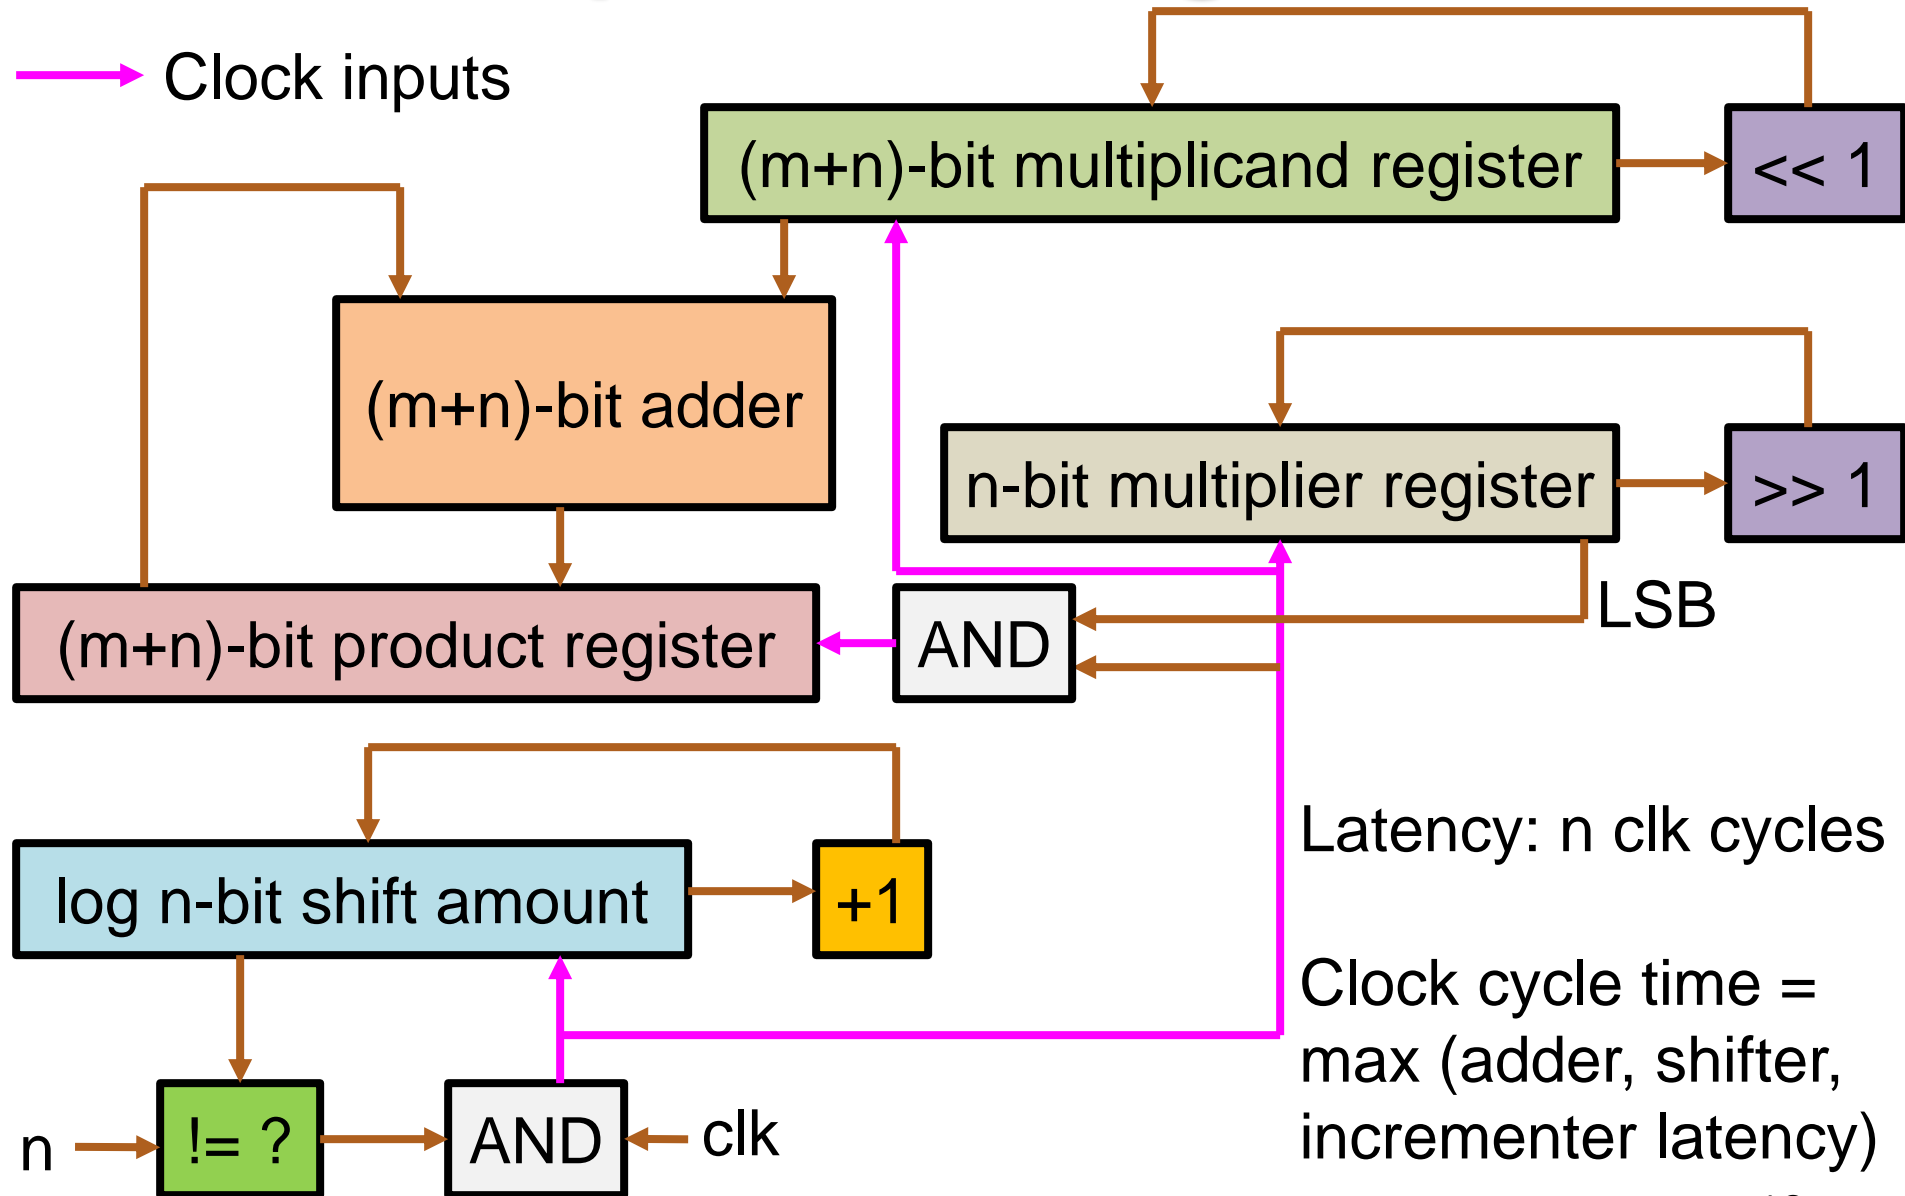
Multiplication algorithm

- Paper-pencil algorithm (unsigned operands)
 - Multiplicand = $(a_{m-1}a_{m-2}\dots a_0) = A$
 - Multiplier = $(b_{n-1}b_{n-2}\dots b_0) = \sum_{i=0}^{n-1} b_i 2^i$
 - Product = $A.b_0 + A.b_1.2 + A.b_2.2^2 + \dots$
 - The left shift operations are needed to take care of the multiplications by 2^i
 - The i^{th} partial product = $(b_i ? A \ll i : 0)$
 - Leads to a simple multiplication algorithm that in each step produces one partial product by inspecting a bit position of the multiplier and accumulates the partial product into a product register
 - We will refer to the i^{th} bit of multiplier as `multiplier[i]`

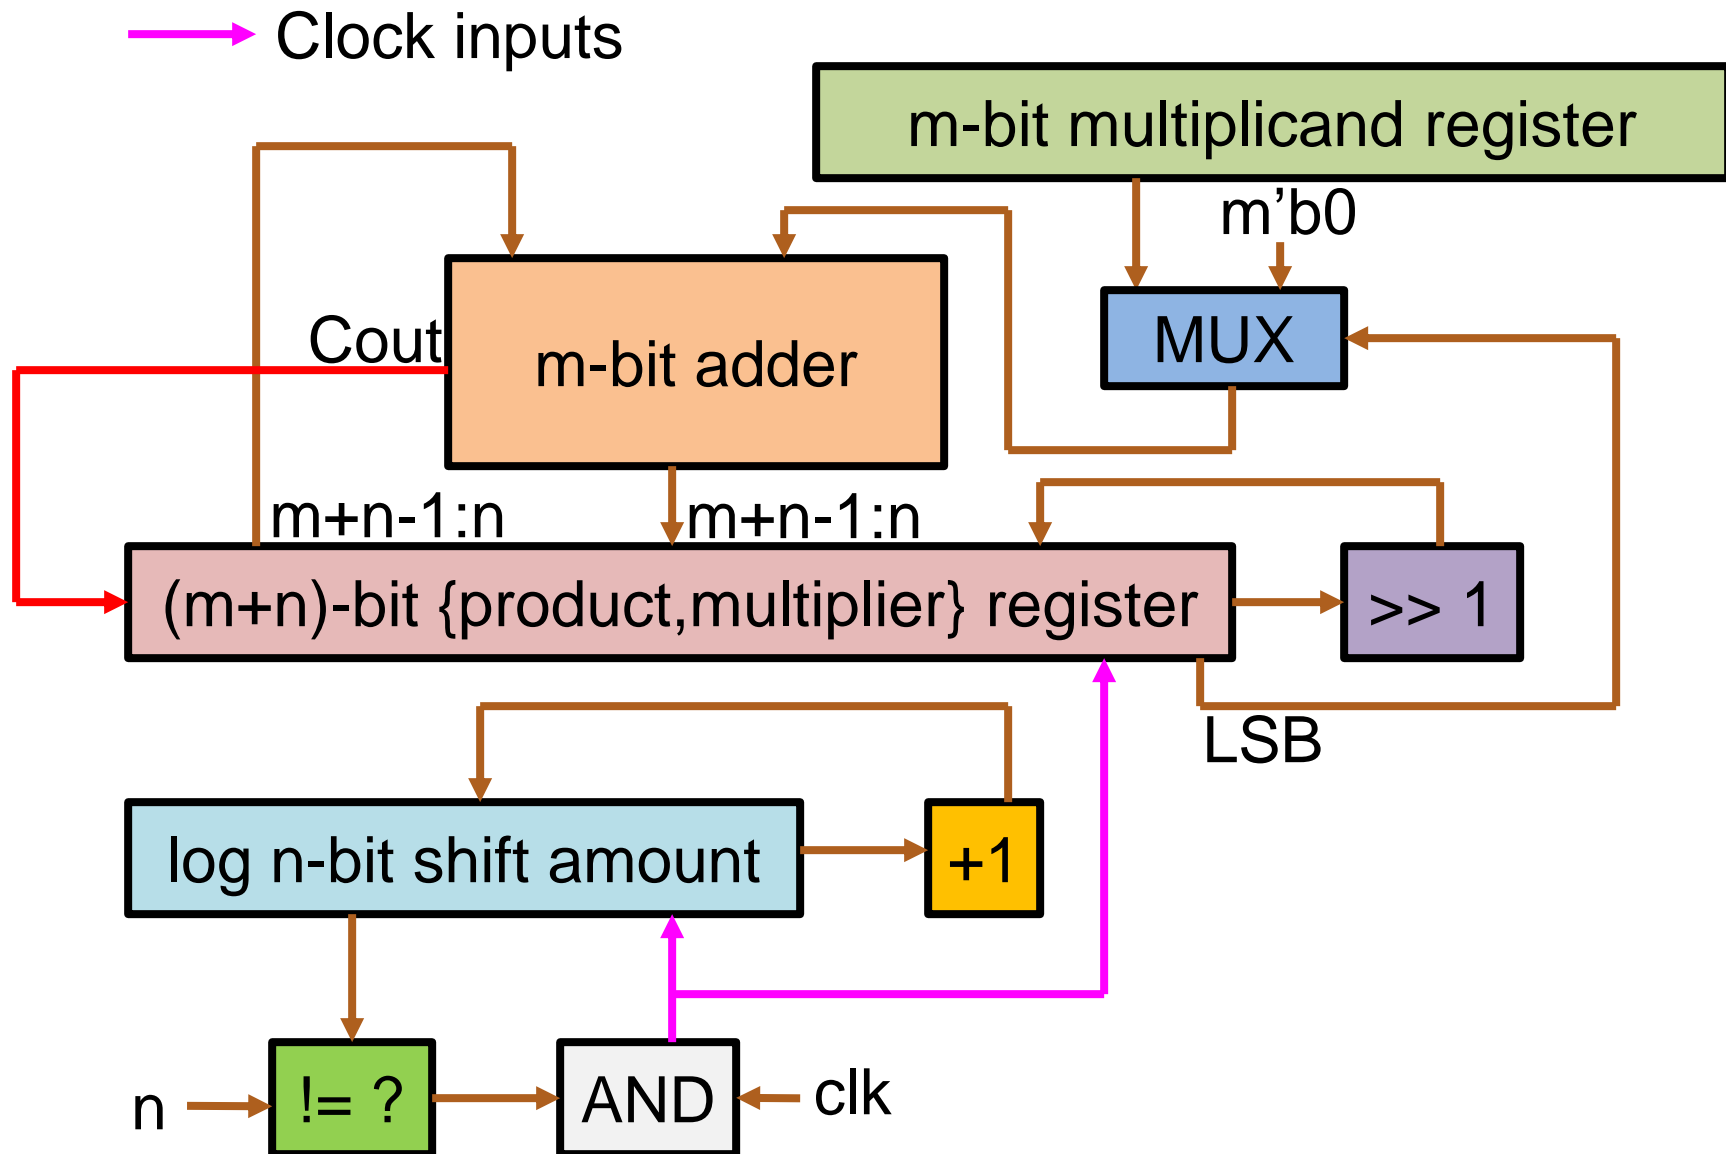
Multiplication algorithm



Multiplication algorithm



Multiplication algorithm (simpler)



Faster multiplication

- In the basic algorithm, n additions are done sequentially because we have used only one adder
- Consider the following grouping
 - Product = $(A.b_0 + A.b_1.2) + (A.b_2.2^2 + A.b_3.2^3) + (A.b_4.2^4 + A.b_5.2^5) + \dots$
 - The additions in the parentheses can be done in parallel after appropriately shifting the multiplicand (requires $n/2$ adders)
 - The results of these $n/2$ additions can be further paired and combined in $n/4$ adders
 - Overall, $n - 1$ additions would be arranged in a binary tree and the multiplication requires $\log_2 n$ cycles (can be done with $n/2$ adders)

Signed multiplication

- Negate the negative operands and do the multiplication
- If the signs of the operands are opposite, negate the multiplication result
- MIPS has both signed and unsigned multiplication
 - mult instruction treats operands as signed
 - multu treats operands as unsigned
- MIPS offers no hardware support for detecting overflow in multiplication
 - If needed, must be done in software

Booth's multiplication algorithm

- Due to Andrew Booth [1951]
- Handles both unsigned and signed operands
- Basic observation
 - Let the multiplicand be A and let the multiplier be $b_{n-1}b_{n-2}\dots b_0 = 000\dots 011\dots 100\dots 0$ where the block 1's ranges from b_p to b_q with $p < q$
 - Therefore, the multiplier is $2^{q+1} - 2^p$ and the product is $(A \ll (q+1)) - (A \ll p)$
 - Can be generalized by representing the multiplier as addition of disjoint blocks of 1's
- Append $b_{-1} = 0$ at the end of the multiplier
- Assume that both multiplicand and multiplier are in two's complement representation

Booth's multiplication algorithm

- Algorithm
 - Initialize product register to 0
 - Scan the multiplier from b_0 to b_{n-1}
 - When examining b_i , if $b_i == b_{i-1}$, keep scanning
 - When examining b_i , if b_i is 1 and b_{i-1} is 0, subtract ($A \ll i$) from product register
 - When examining b_i , if b_i is 0 and b_{i-1} is 1, add ($A \ll i$) to the product register
 - If $i == n - 1$, stop. Final product is in two's complement representation.

Booth's multiplication algorithm

- Observations
 - In the worst case, needs n addition operations when the multiplier has an alternating bit pattern i.e., 0101...01
 - If b_{n-1} is 1 (negative multiplier), the last operation done on the product would be a subtraction
 - Will result in a positive product, if the multiplicand is also negative; otherwise the product is negative
- Booth's algorithm can be implemented in two stages
 - First stage scans the multiplier and prepares the addition/subtraction operands
 - Second stage does the additions and subtractions in a binary tree of adders as in the fast multiplication algorithm

Division algorithm

- Suppose dividend $(A) = (a_{m-1}a_{m-2}\dots a_0)$, divisor $(B) = (b_{n-1}b_{n-2}\dots b_0)$
- Problem statement: find Q and R such that $A=BQ + R$ and $R < B$
 - We will assume unsigned A and B to start with
 - Clearly the maximum length of Q is $m - n + 1$ if $m \geq n$; otherwise Q is 0 and $R = A$
 - Mathematically, we want
$$\sum_{i=0}^{m-1} a_i 2^i = \sum_{i=0}^{n-1} b_i 2^i \sum_{i=0}^{m-n} q_i 2^i + \sum_{i=0}^{n-1} r_i 2^i$$
 - The usual division algorithm proceeds by computing q_i 's starting from the most significant bit i.e., q_{m-n}

Division algorithm

- Steps in usual division algorithm
 - 1. Place A above B such that their most significant bits are aligned
 - This may require shifting B to the left by $m - n$ positions (equivalent to multiplying B by 2^{m-n})
 - E.g., suppose $A = 1001010$, $B = 1000$ ($m=7$, $n=4$)
 $A = 1001010$
 $B' = 1000000$ (shifted B)
 - 2a. If B' is less than or equal to A, subtract B' from A and assign $Q = (Q \ll 1) \mid 1$, $A =$ the result of subtraction, $B' = B' \gg 1$ (to compute next lower q_i ; notice the similarity with polynomial division)
 - 2b. If $B' > A$, shift B' to the right by one bit position and assign $Q = (Q \ll 1)$ [Q gets a 0 bit]
 - Repeat step 2a or 2b $m - n + 1$ times, each time producing one bit of Q; R is residual A at the end

Division algorithm

- Execution of the usual algorithm (iteration 1)

A = 1001010

Q = 0

B' = 1000000

0001010 ← new A

Q = 1

0100000 ← new B'

Division algorithm

- Execution of the usual algorithm (iteration 2)

A = 0001010

Q = 1

B' = 0100000

No subtraction

0001010 \leftarrow A

Q = 10

0010000 \leftarrow new B'

Division algorithm

- Execution of the usual algorithm (iteration 3)

A = 0001010

Q = 10

B' = 0010000

No subtraction

0001010 \leftarrow A

Q = 100

0001000 \leftarrow new B'

Division algorithm

- Execution of the usual algorithm (iteration 4)

$$A = 0001010$$

$$Q = 100$$

$$B' = 0001000$$

$$0000010 \leftarrow \text{new } A$$

$$Q = 1001$$

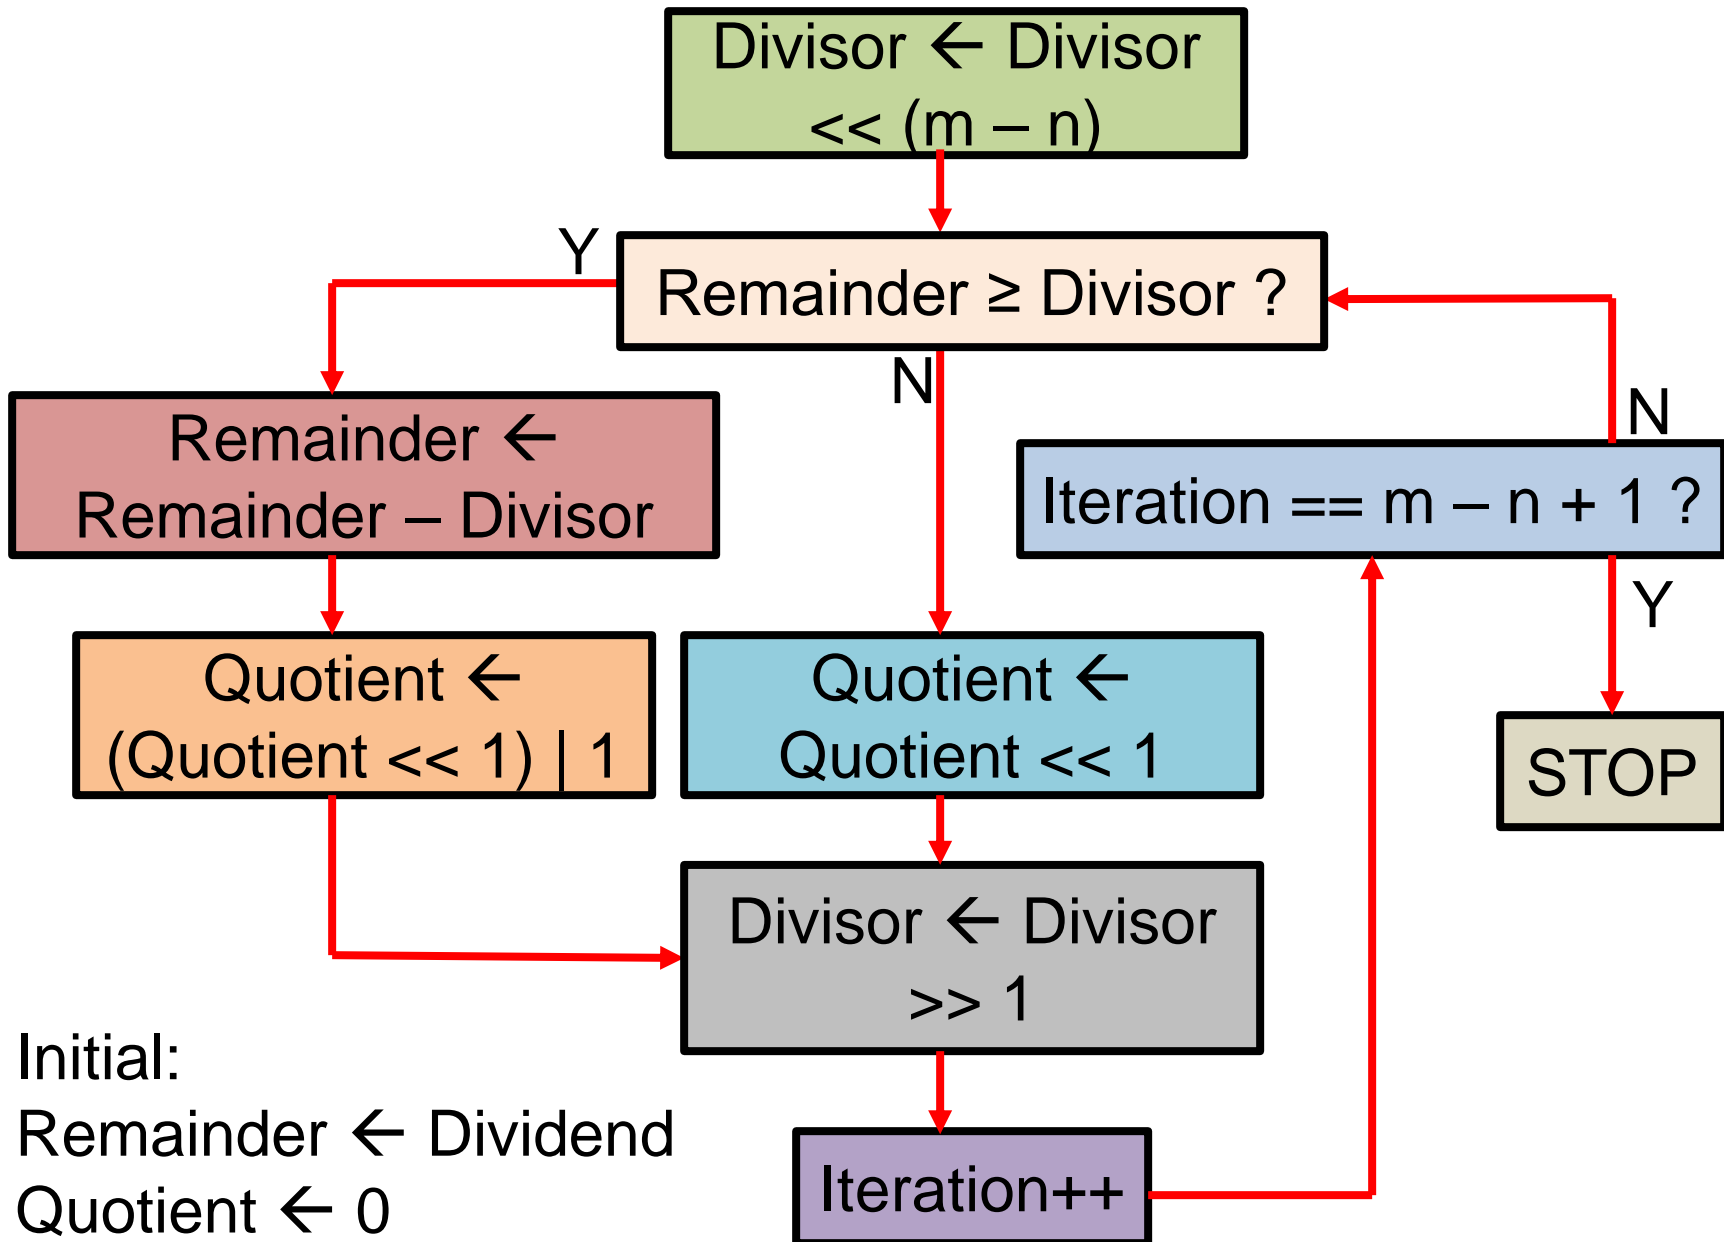
$$0000100 \leftarrow \text{new } B'$$

- Done executing four iterations ($m - n + 1$ iterations)
- Quotient is 1001 and remainder is residual A i.e., 10
- Each iteration computes one coefficient q_i (either 0 or 1) of the polynomial $q_i 2^i$ by equating the leading powers of residual A and B

Division algorithm

- The first step of aligning the divisor and dividend is important for the main division algorithm to work correctly
 - In a processor, A and B will be in two 32-bit registers to begin with
 - The first step requires computing $m - n$ and shifting B to the left by $m - n$ bits
 - How to compute m and n quickly from A and B?
 - This is essentially a search problem that searches for the most significant non-zero bit in a given binary string

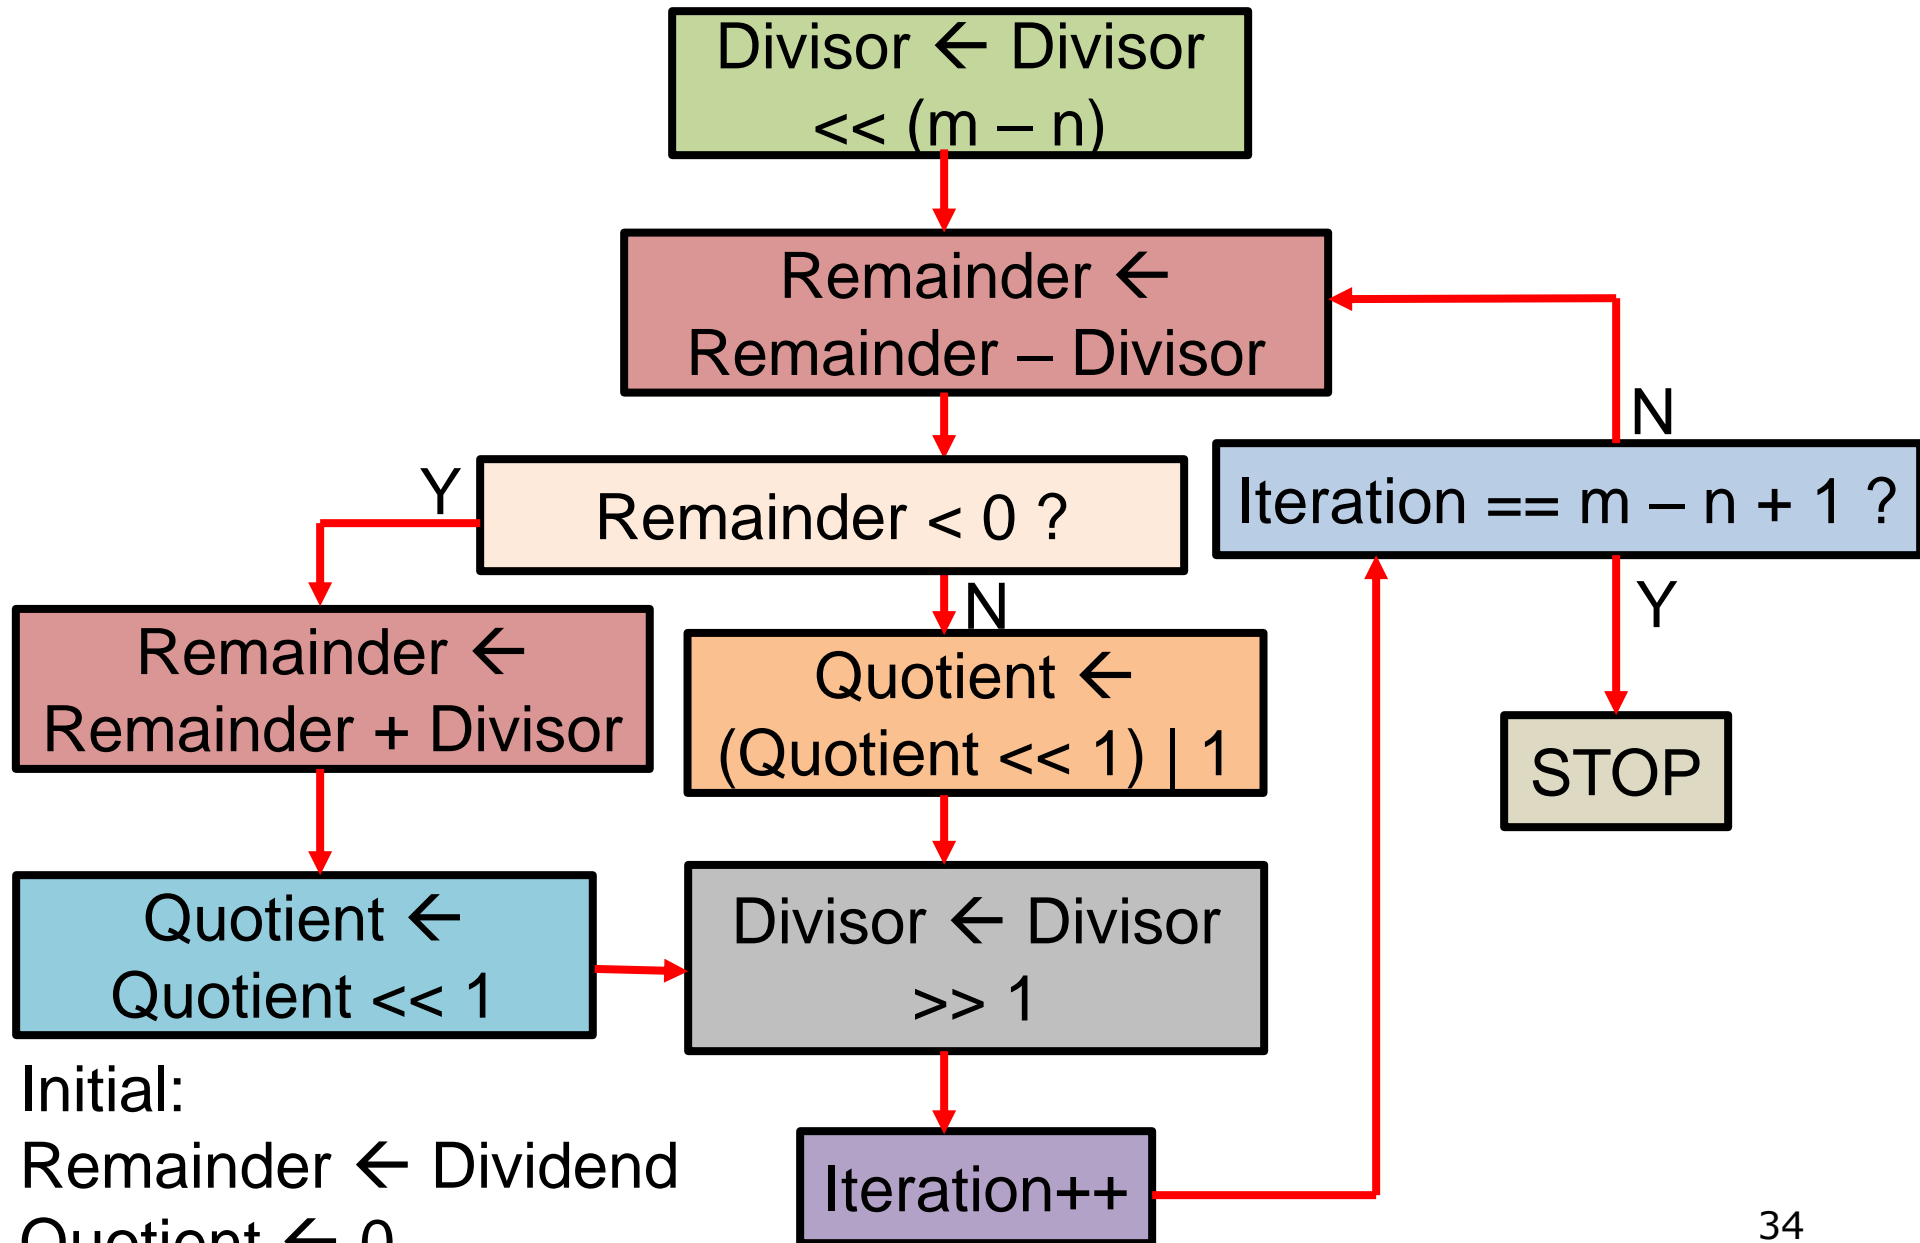
Division algorithm



Division algorithm

- Critical path of one iteration involves a comparison and a subtraction
 - Left shift and OR, right shift, increment can all be done in parallel with subtraction
- A variant of this algorithm always subtracts, then checks if the remainder is negative and if yes, restores the remainder
 - Known as restoring division algorithm
 - Checking if remainder is negative requires comparing only the most significant bit
 - Critical path now involves a subtraction, a negative check, and an addition

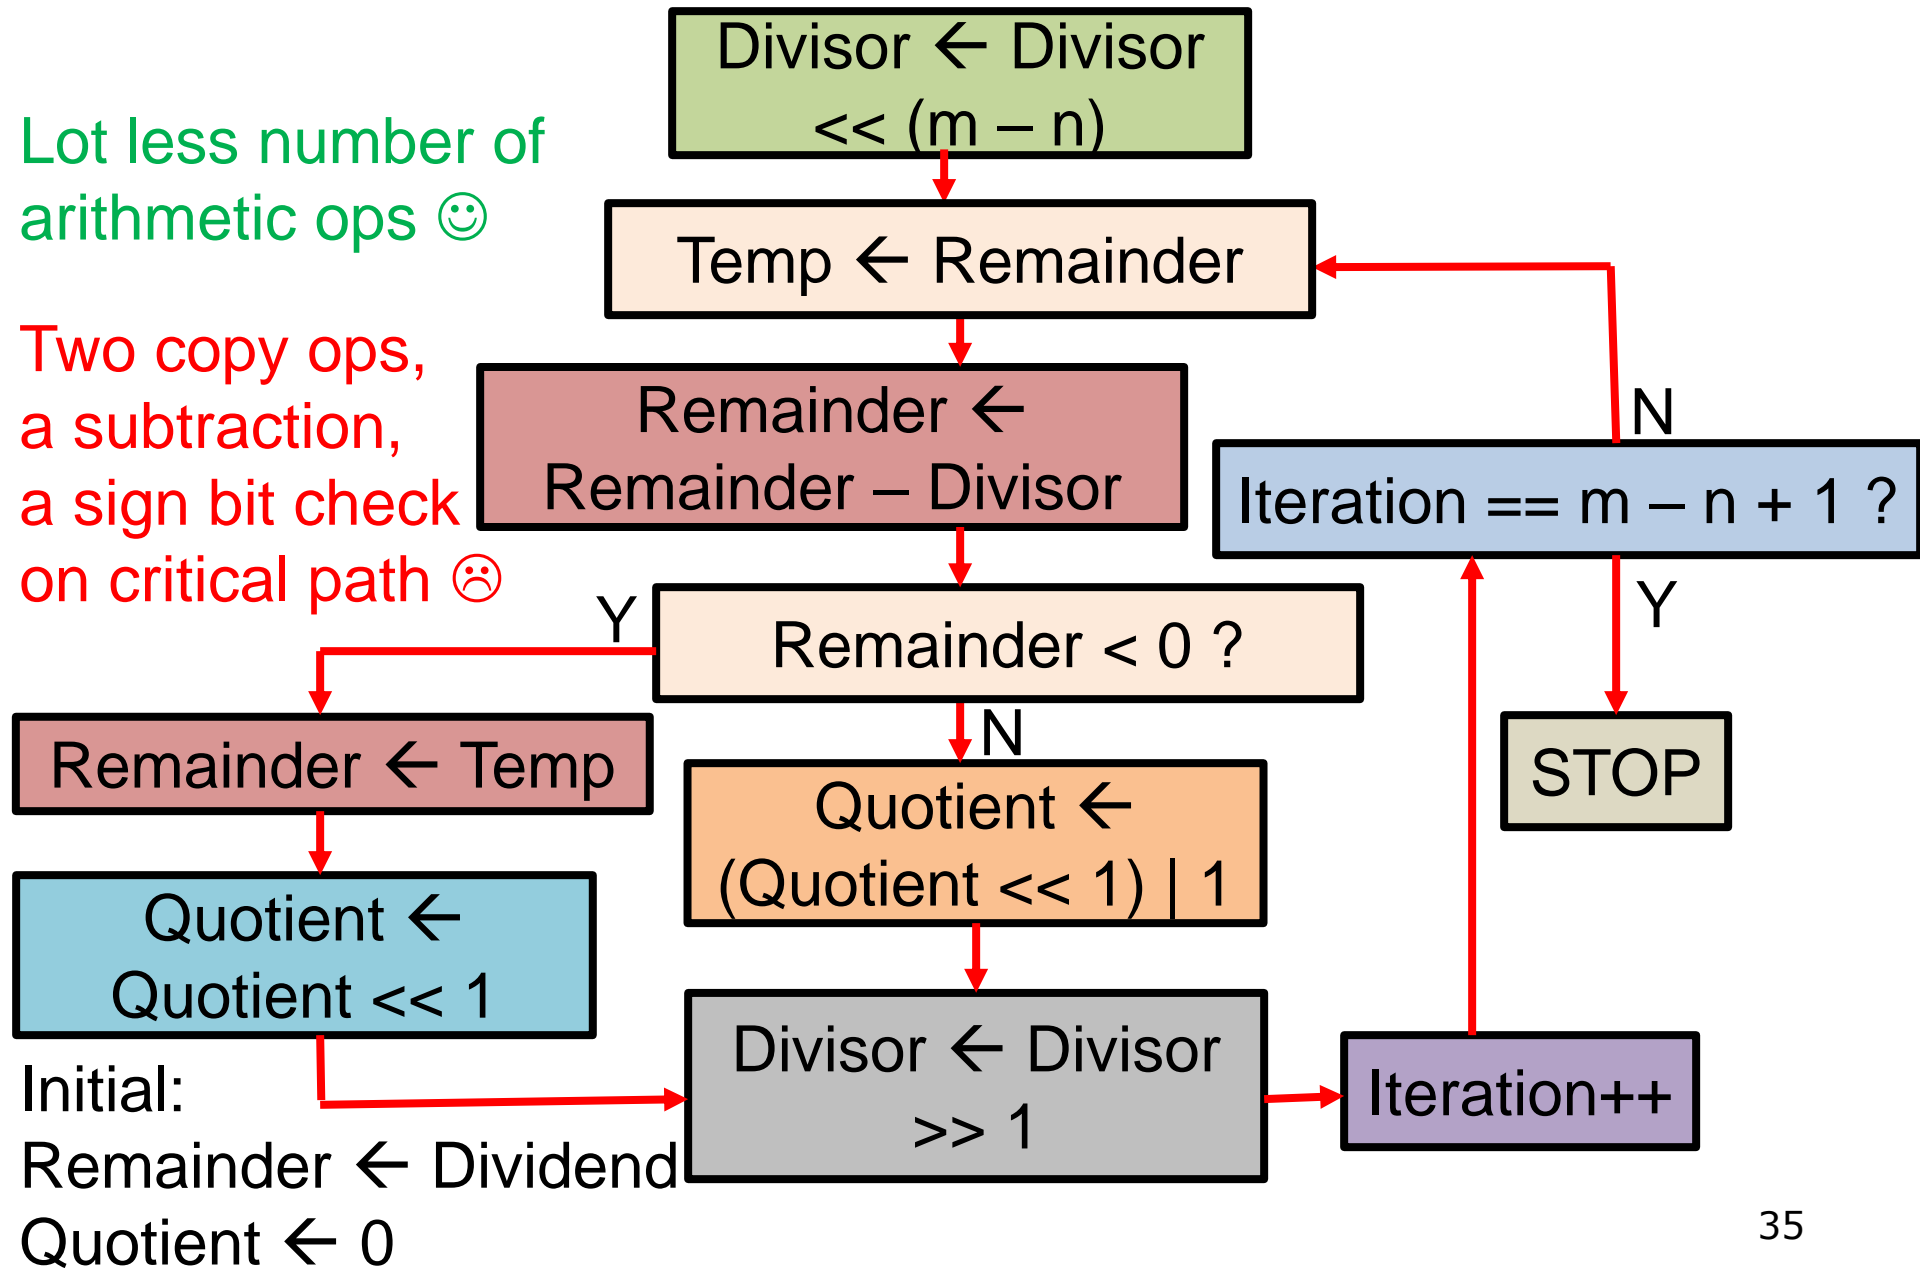
Restoring division algorithm



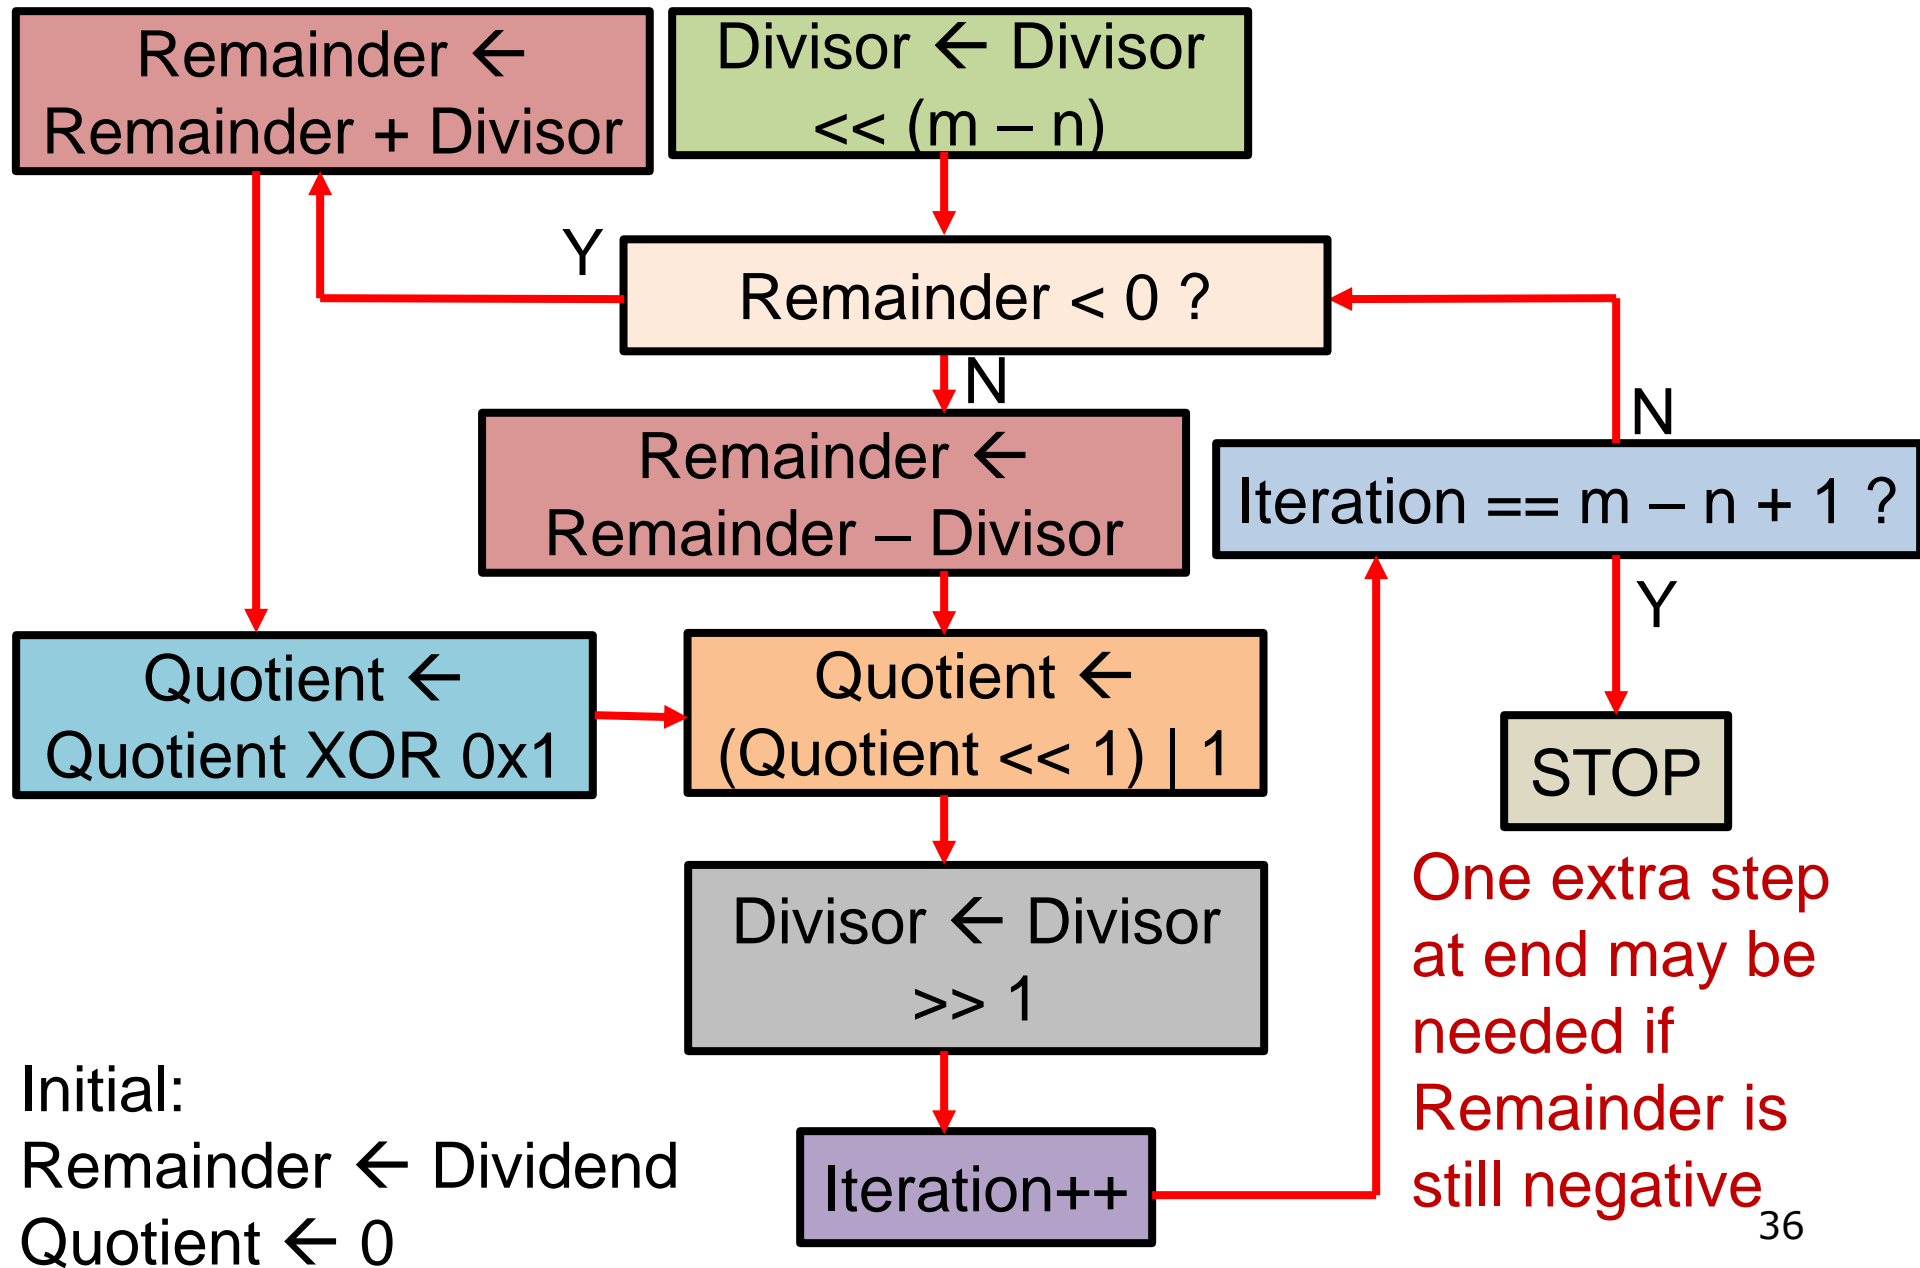
Non-performing restoring division

Lot less number of arithmetic ops 😊

Two copy ops, a subtraction, a sign bit check on critical path ☹️



Non-restoring division algorithm



Non-restoring division algorithm

- Why is this correct?
 - Suppose in iteration#k of restoring division, $r - d$ turns out to be negative, so we restore r
 - In iteration#k+1 of restoring division, we compute $r - d/2$ (note that d is shifted right)
 - In non-restoring division's iteration#k, we keep $r - d$ (which is negative) and in iteration#k+1, we do $r - d + d/2$ yielding $r - d/2$, which is same as the value of the remainder in restoring division after iteration#k+1
 - The quotient is adjusted in iteration#k+1 by first making the last shifted bit 0 and then shifting in a 1
- Critical path = sign bit check; (addition or subtraction || Quotient update || divisor shift || iteration increment) [";" \rightarrow seq., "||" \rightarrow parallel]

Non-restoring division algorithm

- Hardware implementation
 - Directly follows from the algorithm
 - It is possible to combine {remainder, quotient} in a single register of size m bits
 - At the beginning, the entire register is occupied by the m -bit dividend and quotient is zero
 - At the end, the least significant $m - n$ bits along with the last quotient bit will be the final quotient and the remaining n bits will contain the remainder
 - Justifies why Lo contains quotient and Hi contains remainder in MIPS division instructions' result
 - Key observation: instead of shifting the divisor to the right in each iteration, the same effect can be achieved by shifting the dividend to the left

Division algorithm (A shifted left)

- Execution of the usual algorithm (iteration 1)

$A, Q = 1001010$

$Q = 0$

$B' = 1000000$

0001010

$Q = 1$

0010101 \leftarrow new A, Q

Division algorithm (A shifted left)

- Execution of the usual algorithm (iteration 2)

$A, Q = 0010101$

$Q = 1$

$B' = 1000000$

No subtraction

$0101010 \leftarrow \text{new } A, Q$

$Q = 10$

Division algorithm (A shifted left)

- Execution of the usual algorithm (iteration 3)

$A, Q = 0101010$

$Q = 10$

$B' = 1000000$

No subtraction

$1010100 \leftarrow \text{new } A, Q$

$Q = 100$

Division algorithm (A shifted left)

- Execution of the usual algorithm (iteration 4)

A, Q = 1010100

Q = 100

B' = 1000000

0010100

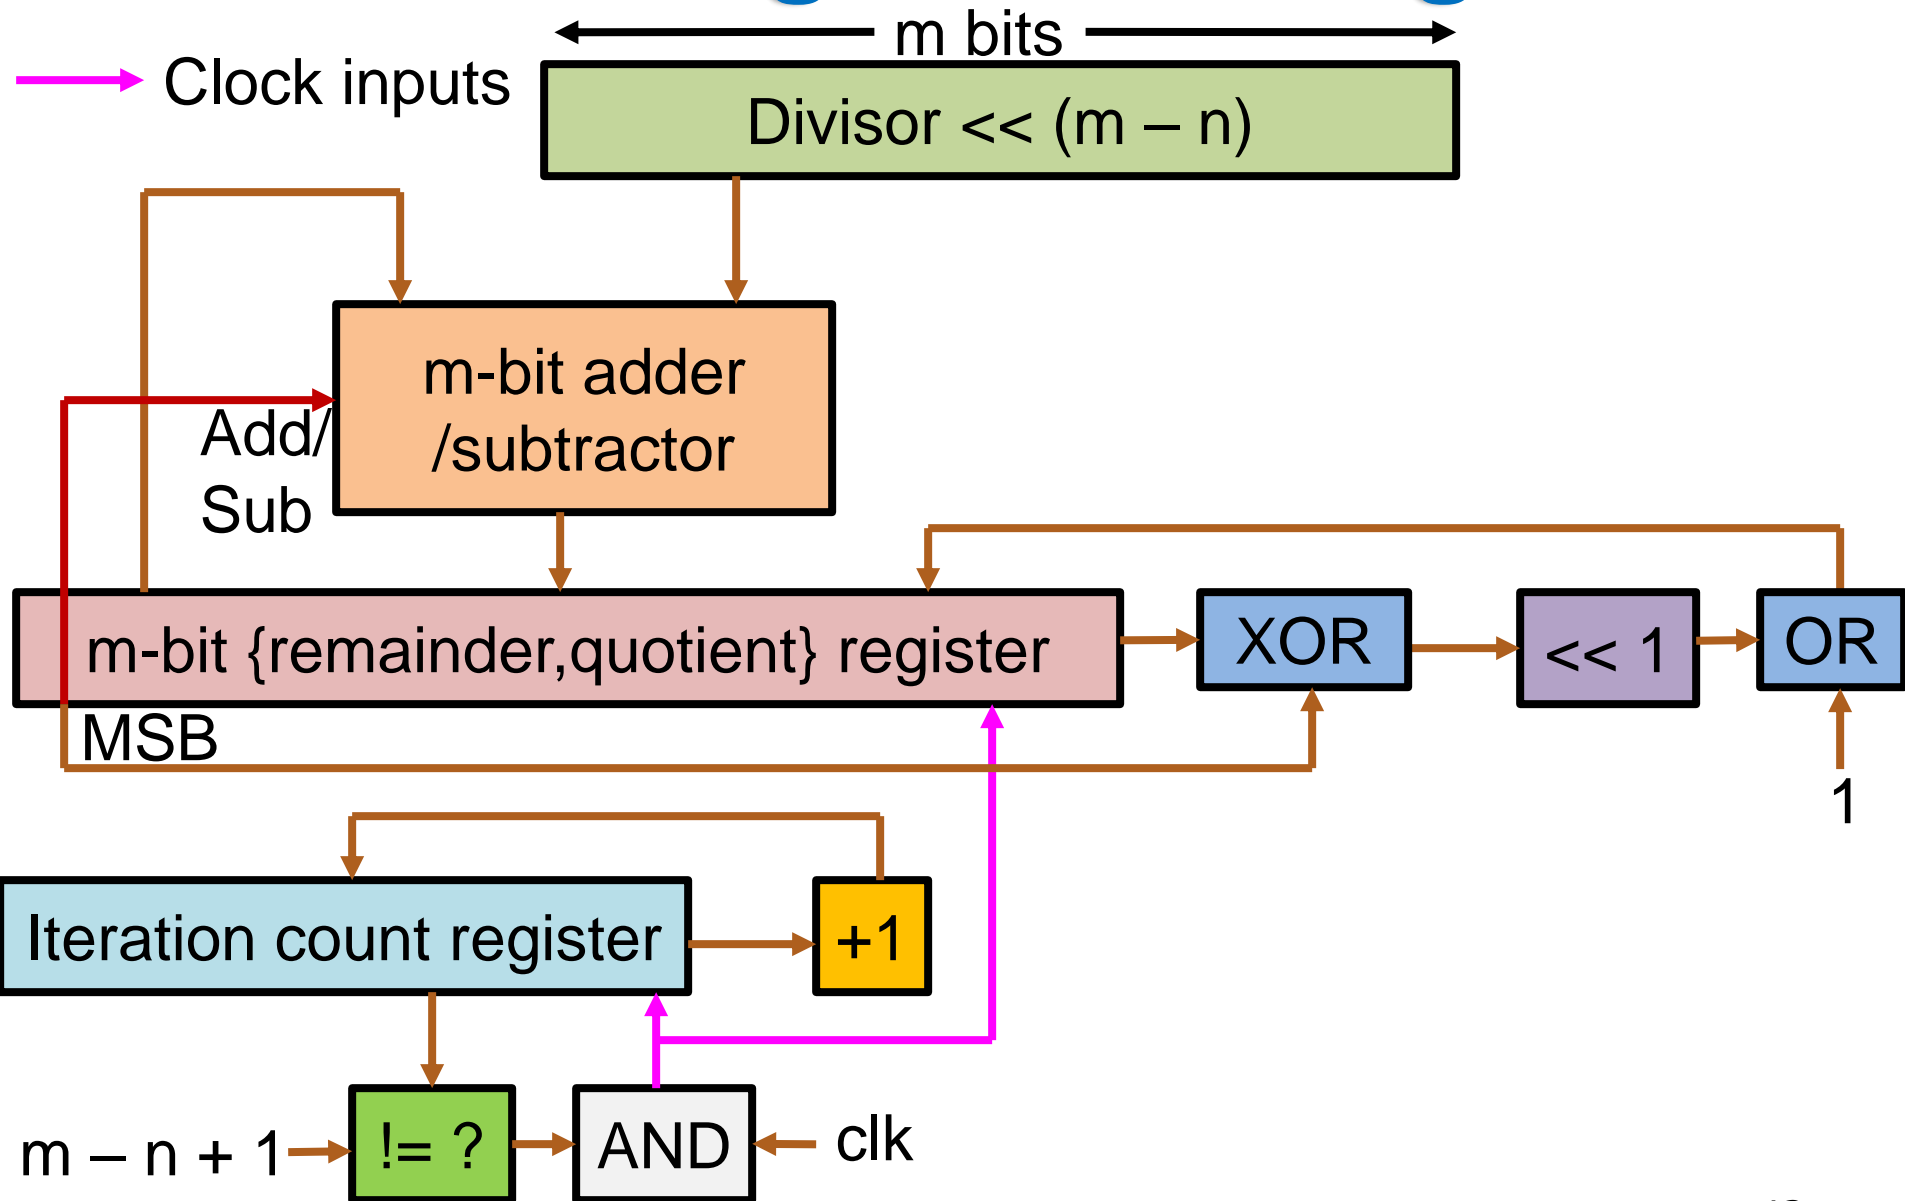
Q = 1001

Remainder

Quotient = 1001

- Done executing four iterations ($m - n + 1$ iterations)
- The remainder appears in upper 4 bits ($n=4$) of residual $\{A, Q\}$ and quotient is the concatenation of the lower 3 bits ($m - n = 3$) of residual $\{A, Q\}$ and the last bit of quotient

Non-restoring division algorithm



Signed division

- Suppose A and B are positive and we have $A = BQ + R$
- So, $-A = -BQ - R$ is the expected result of dividing $-A$ by B
- Rule: remainder must have the same sign as the dividend

Signed division

- Signed division algorithm
 - Convert the negative operand(s) to their positive value(s) and use any unsigned division algorithm
 - If both operands were positive originally, stop
 - $A = BQ + R$
 - If both operands were negative originally, negate remainder
 - $-A = -BQ + (-R)$
 - If dividend was negative and divisor was positive, negate both quotient and remainder
 - $-A = B(-Q) + (-R)$
 - If dividend was positive and divisor was negative, negate quotient
 - $A = -B(-Q) + R$

Signed division

- MIPS has both signed and unsigned division
 - div instruction treats operands as signed
 - divu treats operands as unsigned
- Faster division algorithms are non-trivial
 - The subtractions cannot be done in parallel
 - Fast algorithms try to predict quotient bits and detect and correct mispredictions on the fly

Floating-point addition

- Assume that the operands are represented in IEEE 754 format
- Step 1: Align exponents of the operands by shifting significand of the smaller number to the right (recall significand = 1.mantissa)
- Step 2: Add significands
 - Use 2's complement integer arithmetic ignoring the binary point
 - If result is negative, convert to its positive value
 - Put back the binary point after addition and assign the sign bit depending on the result's sign
- Step 3: Normalize the sum
 - Shift right and increase exponent OR shift left and decrease exponent

Floating-point addition

- Step 4: If overflow or underflow detected, raise exception
- Step 5: Round mantissa of the normalized sum to fit the precision
 - This may make the result unnormalized e.g., 1.11111...1 in binary becomes 10.00000...0 after half-way rounding rule is applied
- If not normalized, repeat steps 3, 4, 5; else stop
- Precision could be lost in intermediate computations e.g., during shifting significand
 - To avoid this, IEEE 754 standard recommends use of three additional bits in intermediate representations (guard bit, round bit, sticky bit)

Use of guard, round, sticky bits

- The guard and round bits extend the mantissa by two bits
- The sticky bit appears after the round bit and is set to 1 if any bit to the right of the round bit is 1
- These bits improve the precision of arithmetic whenever the significand needs to be shifted to the right
 - These bits retain a few mantissa bits which would otherwise get dropped
- If the normalization step needs to shift the significand to the left, these bits will start contributing to the precision of the mantissa

Floating-point addition

- Example: $0.5 + (-0.4375)$
 - $0.5 = 1.000 \times 2^{-1}$, $-0.4375 = -1.110 \times 2^{-2}$
- Step 1: Align exponents
 - $-1.110 \times 2^{-2} = -0.111 \times 2^{-1}$
- Step 2: Add significands
 - $1.000 + (-0.111)$
 - $1000 + 1001 = 0001$ ($2'$ complement)
 - 0.001 after putting back the binary point
- Step 3: Normalize
 - $0.001 \times 2^{-1} = 1.000 \times 2^{-4} = 0.0625$
 - No overflow or underflow because exponent is within legitimate range; no need to round also

Floating-point multiplication

- Step 1: Add the biased exponents of the operands and subtract the bias to get the new biased exponent
- Step 2: Multiply significands
 - Use unsigned integer multiplication algorithm ignoring the binary point
 - Put the binary point after N bit positions from right where N is the sum of the lengths of the mantissa of two operands
- Step 3: Normalize the product
- Step 4: If overflow or underflow detected, raise exception
- Step 5: Round the product; if product is not normalized, repeat steps 3, 4, 5

Floating-point multiplication

- Step 6: If the signs of the operands are same, the product is assigned positive sign; otherwise the product is negative
- Example: $0.5 \times (-0.4375)$
 - $(1.000 \times 2^{-1}) \times (-1.110 \times 2^{-2})$
 - New biased exponent = $(127 - 1) + (127 - 2) - 127 = 124 = (127 - 3)$
 - Actual exponent -3
 - Multiply significands: $1000 \times 1110 = 1110000$
 - Put back binary point: 1.110000
 - Unsigned product = 1.110×2^{-3}
 - Already normalized and no overflow
 - No need to round
 - Final product = $-1.110 \times 2^{-3} = -0.21875$