

Processor FSM Design

Mainak Chaudhuri
Indian Institute of Technology Kanpur

Sketch

- Abstract model of computer
- Single-cycle instruction execution
- Multi-cycle instruction execution
- Pipelined instruction execution

2

Abstract model of computer

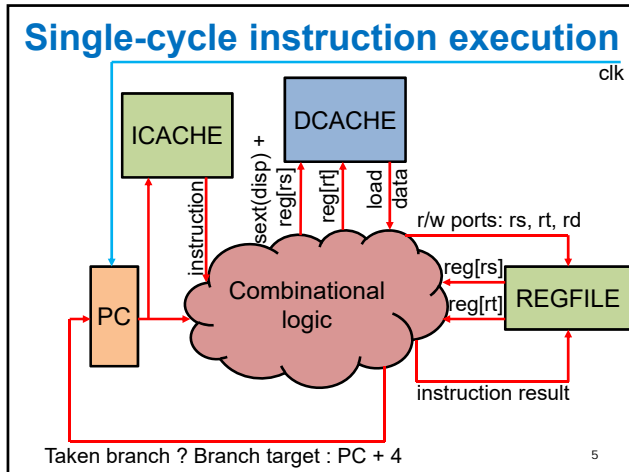
- Each instruction undergoes five stages
 - Stage 0 (IF): fetch the instruction pointed to by program counter from memory
 - Stage 1 (ID/RF): decode the instruction to extract various fields and read source register operands
 - Source registers can be read in parallel with instruction decoding in MIPS due to fixed positions of rs and rt in all formats of the ISA
 - Stage 2 (EX): execute the instruction in ALU; compute address of load/store instructions; update program counter to PC+4 or branch target (if this instruction is control transfer ins)
 - Stage 3 (MEM): access memory if load/store instruction; use address computed in stage 2
 - Stage 4 (WB): write result back to destination register if the instruction produces a result

3

Single-cycle instruction execution

- Possible to implement all five stages as a big combinational logic
 - All the memory elements are also combinational meaning that the input address can be presented any time within a cycle and the output data is available within the same cycle
 - Register file (in stages 1 and 4), icache (in stage 0), and dcache (in stage 3) are the combinational memory elements
 - PC, icache[PC], reg[rs], reg[rt], dcache[disp(rs)] are inputs
 - dcache[disp(rs)] is a relevant input for loads
 - PC, reg[rd] or reg[rt], dcache[disp(rs)] are outputs
 - dcache[disp(rs)] is a relevant output for stores
 - The instruction opcode and function (for R format) fields decide the control of combinational logic e.g., which register operand(s) is/are valid sources, which ALU operation to invoke, etc.

4



Single-cycle instruction execution

- Taken branch

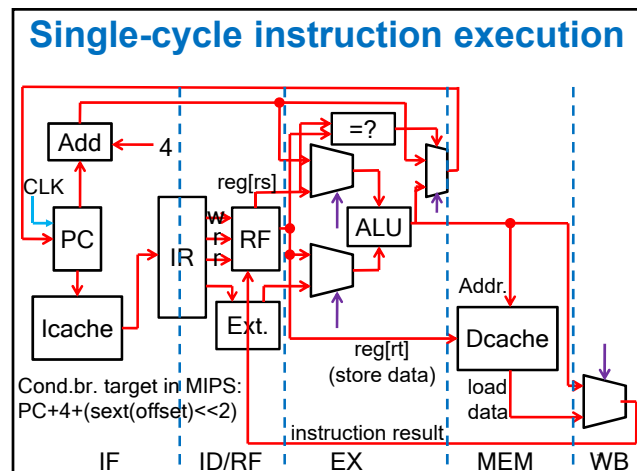
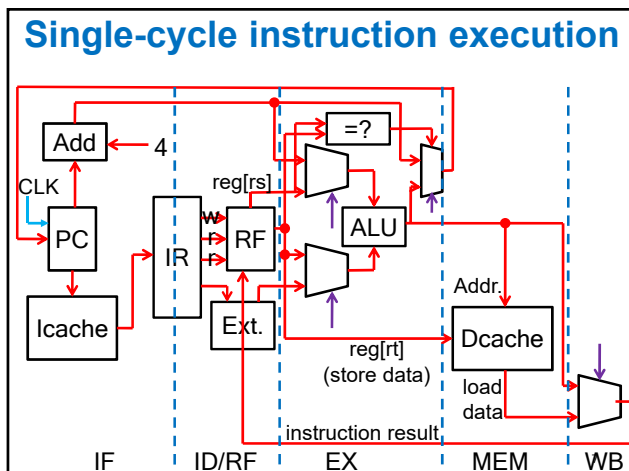
- A branch instruction whose condition dictates that the next instruction to be executed is the one at the branch target

- beq \$1, \$2, label (the branch is taken if $\$1 == \2)
- Unconditional direct and indirect jumps and procedure calls are always taken branches

- Not taken branch

- A branch instruction whose condition dictates that the next instruction to be executed is the instruction right next to the branch

- A la Frost's "Road Not Taken"
- beq \$1, \$2, label (the branch is not taken if $\$1 != \2)
- A conditional branch can be taken or not taken depending on how the condition evaluates at run time



Single-cycle instruction execution

- Clock cycle time must be large enough to accommodate the lengthiest instruction
 - This is typically the load instruction whose critical path includes all five stages: IF: fetch instruction, ID/RF: decode while reading reg[rs] and sign extending displacement, EX: add reg[rs] with sext(displacement) to generate address, MEM: read dcache to get load data, WB: write load data back to reg[rt]
 - Not a good idea since many other instructions take smaller amount of time
 - Control transfer instructions require only three stages
 - Stores require only four stages
 - All ALU instructions can bypass MEM stage

8

Multi-cycle instruction execution

- Each stage takes one cycle to execute
 - Processor is now a five state FSM and the memory elements are sequential
 - Need flip-flops at stage boundaries and each stage is a combinational logic
 - Clock cycle time = latency of the longest stage
 - Most instructions take five cycles
 - ALU instructions do nothing in MEM stage, but will have to go through it if FSM state transition is implemented using an incrementer
 - Some instructions may take longer e.g., multiply/divide, load/store (in case of dcache miss)
 - Some instructions may take shorter time e.g., stores take four cycles, control transfer instructions take three cycles
 - Overall CPI depends on instruction mix

9

Multi-cycle instruction execution

- Example
 - IF: 2 ns, ID/RF: 1 ns, EX: 1 ns, DMEM: 3 ns, WB: 1 ns
 - Branch frequency: 20%
 - Store frequency: 10%
 - Multiply/divide frequency: 5%, latency: 30 ns
 - Total instruction count: 100
- Multi-cycle
 - Cycle time: 3 ns, frequency: 333 MHz
 - CPI: $0.2*3 + 0.1*4 + 0.05*10 + 0.65*5 = 4.75$
 - Execution time: $100*4.75*3 \text{ ns} = 1425 \text{ ns}$
- Single-cycle
 - Cycle time: 8 ns, frequency: 125 MHz
 - CPI: $1*0.95 + \text{ceil}(30/8)*0.05 = 1.15$
 - Execution time: $100*1.15*8 \text{ ns} = 920 \text{ ns}$

10

Multi-cycle instruction execution

- Same example with a balanced pipeline
 - IF: 2 ns, ID/RF: 2 ns, EX: 2 ns, DMEM: 2 ns, WB: 2 ns
 - Branch frequency: 20%
 - Store frequency: 10%
 - Multiply/divide frequency: 5%, latency: 30 ns
 - Total instruction count: 100
- Multi-cycle
 - Cycle time: 2 ns, frequency: 500 MHz
 - CPI: $0.2*3 + 0.1*4 + 0.05*15 + 0.65*5 = 5$
 - Execution time: $100*5*2 \text{ ns} = 1000 \text{ ns}$
- Single-cycle
 - Cycle time: 10 ns, frequency: 100 MHz
 - CPI: $1*0.95 + \text{ceil}(30/10)*0.05 = 1.1$
 - Execution time: $100*1.1*10 \text{ ns} = 1100 \text{ ns}$ (worse than multi-cycle)

11

Pipelined instruction execution

- Observations from multi-cycle design
 - In the second cycle, I know if it is a branch; if not, start fetching the next instruction?
 - When the ALU is doing an addition (say), the decoder is sitting idle; can we use it for some other instruction?
 - In summary, exactly one stage is active at any point in time: wastes hardware resources
- Form a pipeline
 - Process five instructions in parallel
 - Each instruction is in a different stage of processing (called pipe stage)

12

Pipelined instruction execution

- Individual instruction latency is five cycles, but ideally can finish one instruction every cycle after the pipeline is filled up
 - Ideal CPI of 1.0 at the clock frequency of multi-cycle design
 - Execution time is ideally one-fifth of the multi-cycle design
 - Instruction throughput improves five times (number of instructions completed in a given time)

I0	IF	ID/RF	EX	MEM	WB				
I1		IF	ID/RF	EX	MEM	WB			
I2			IF	ID/RF	EX	MEM	WB		
I3				IF	ID/RF	EX	MEM	WB	
I4					IF	ID/RF	EX	MEM	

→ time

Pipelined instruction execution

- Gains
 - Extracting parallelism from a sequential instruction stream: known as instruction-level parallelism (ILP)
 - Can complete one instruction every cycle (ideally)
- Loss
 - Each pipe stage may get lengthened a little bit due to control overhead (skew time, setup time, propagation delay): limits the gain due to pipelining
 - Each instruction may take slightly longer for this reason
 - Bigger aggregate memory bandwidth: icache and dcache may miss in the same cycle
 - Pipeline hazards

14


Pipeline hazards: structural

- Structural hazards
 - Arises due to resource conflicts
 - Happens if the same resource is accessed in at least two stages of the pipe
 - Fewer resources than needed
 - Unpipelined functional units
 - Lack of memory or register file ports
- Why fewer resources?
 - Reduction in complexity (and power consumption)
 - Make the common case fast: pipelined divider may only waste silicon estate
- All types of hazards introduce stalls or pipeline bubbles

15

Pipeline hazard: control

- Branches pose a problem




- Two pipeline **bubbles**: increases average CPI
- Can we reduce it to one bubble?

16

Branch delay slot

- MIPS R3000 has one bubble
 - Called branch delay slot
 - Exploit clock cycle phases
 - On the positive half compute branch condition
 - On the negative half fetch the target



- The PC update hardware (selection between target and next PC) works on the lower edge

Branch delay slot

- Can we utilize the branch delay slot?
 - The delay slot is always executed (irrespective of the fate of the branch)
 - Reason why branch target is $PC+4+(\text{sext}(\text{offset})\ll 2)$
 - Boost instructions common to fall through and target paths to the delay slot or from earlier than the branch
- ```

if (cond) {
 // fall through
}
else {
 // target
}

```

18

## Branch delay slot

- Can we utilize the branch delay slot?
    - The delay slot is always executed (irrespective of the fate of the branch)
      - Reason why branch target is  $PC+4+(\text{sext}(\text{offset})\ll 2)$
    - Boost instructions common to fall through and target paths to the delay slot or from earlier than the branch
- ```

if (cond) {
    // fall through
}
else {
    // target
}
  
```

18

Branch delay slot

- Can we utilize the branch delay slot?
 - The delay slot is always executed (irrespective of the fate of the branch)
 - Reason why branch target is $PC+4+(\text{sext}(\text{offset})\ll 2)$
 - Boost instructions common to fall through and target paths to the delay slot or from earlier than the branch

```

if (cond) {
    // fall through
}
else {
    // target
}

```

branch instruction

delay slot

18

Branch delay slot

- Can we utilize the branch delay slot?
 - Boost instructions common to fall through and target paths to the delay slot or from earlier than the branch
 - Not always possible to find
 - Must boost something that does not alter the outcome of the computation
 - If the BD slot is filled with useful instruction then we don't lose anything in CPI; otherwise we pay a **branch penalty** of one cycle

19

Branch prediction

- To increase the clock frequency, designers may choose a deeper pipeline
 - Subdivide the longest stage further
- If the number of pipe stages increases between IF and EX, branch penalty also increases
 - Pipe1: IF ID/RF EX ... Br. penalty?
 - Pipe2: IF ID RF EX ... Br. penalty?
 - Pipe3: IF1 IF2 ID RF1 RF2 EX ... Br. penalty?
 - Assume no branch delay slot
 - Meaning that IF or EX cannot be accommodated in half cycle

20

Branch prediction

- To increase the clock frequency, designers may choose a deeper pipeline
 - Subdivide the longest stage further
- If the number of pipe stages increases between IF and EX, branch penalty also increases
- Today all processors rely on branch predictors that observe the behavior of individual branches and learn to predict their future behavior
 - Makes it possible to infer the next instruction's PC even before the branch executes
 - Pipeline must be flushed on a wrong prediction

Branch target buffer

- Branch target buffer
 - We can put a branch target cache in the fetcher
 - Also called branch target buffer (BTB)
 - Use the lower bits of the instruction PC to index the BTB
 - Use the remaining bits to match the tag
 - In case of a hit the BTB tells you the target of the branch when it executed last time
 - You can hope that this is correct and start fetching from that predicted target provided by the BTB
 - Later you get the real target, compare with the predicted target, and throw away the fetched instruction in case of misprediction; keep going if predicted correctly

22

Branch target buffer

- BTB is looked up with the PC of every instruction in parallel with fetching the instruction
 - On a hit, it provides two pieces of information: this instruction is a control transfer instruction and the target of this control transfer instruction seen last time
 - This target will be used to fetch in the next cycle
 - On a miss, the fetcher has no option but to fetch from the fall through path (PC+4) in the next cycle
 - A control transfer instruction is inserted in the BTB after the EX stage when its target is known
 - A lookup at this point may hit in the BTB; if the branch is not taken, the BTB entry is invalidated; otherwise the entry is updated with the taken target
 - If a lookup at this point misses in the BTB, a new entry is allocated provided the branch is taken

Branch target buffer

- Assume for a program
 - 90% of all control transfer instructions hit in the BTB
 - 90% of outcomes provided by BTB hits are correct
 - 20% of control transfer instructions that miss in the BTB result in taken branches
 - Fraction of bubbles saved = BTB prediction accuracy = $0.9 \cdot 0.9 + 0.1 \cdot 0.8 = 0.89$
 - 11% branches suffer from mispredictions and will require some "recovery" mechanism for correct execution

24

Branch prediction

- BTB will work great for
 - Loop branches (how many misprediction?)
 - Subroutine calls
 - Unconditional branches
- What about *jalr*?
- Conditional branch prediction
 - Rather dynamic in nature
 - The last target is not very helpful in general (if-then-else)
 - Need a direction predictor (predicts taken or not taken)
 - Once that prediction is available we can compute the target

Direction predictors

- How about static prediction?
 - Always not-taken (NT) or always taken (T): penalty?
 - Forward not-taken and backward taken (rationale?)
- Deeper pipelines
 - Define branch penalty
 - **Big problem: deeper pipelines increase branch penalty**
 - Must have better branch predictors for deeper pipeline

26

Control dependence

- Roughly every fifth instruction is a branch
 - Need to be on the right *control flow path*
 - This is the source of input to the pipeline
 - Static techniques are not enough
 - Need highly accurate dynamic predictors
 - Speculate past branches: Alpha 21264 allows 20 outstanding branches, MIPS R10000 allows only 4
 - Need to speculate past predicted branches in deeper pipelines (not a big issue in five-stage pipe)
 - Prediction accuracy?
 - Probability of a correct prediction is p
 - Probability of staying on correct path after n predictions p^n
 - What is minimum p if n is 4 (MIPS R10k)? If n is 20?

Direction predictors

- Let us encode each taken branch as 1 and not taken branch as 0
 - The behavior of a conditional branch can be represented as a binary string
 - Loop branch: 11111...110
 - Alternating if-else: 1010101010 or 0101010101
 - If-else branches can exhibit a wide variety of patterns

Direction predictors


- The problem of direction prediction is essentially design of an estimator that, given an n -bit history, tells us the next most likely outcome for a particular static branch instruction
 - All branch predictors compute the probability of seeing a zero or one, given the recent pattern history h of some limited length n
 - Suppose the number of time 0 appears after h is C_0 and similarly define C_1 ; the prediction is 1 if $C_1 \geq C_0$ and 0 otherwise
 - Instead of having two counters, $C_1 - C_0$ is maintained

Direction predictors

- Let the difference counter be $C(h)$ for a certain history h : $C(h) = C1(h) - C0(h)$
 - Let $C(h)$ be of k bits length (this is independent of the history length)
 - Can count from 0 to $2^k - 1$
 - On seeing 0 after h , decrement $C(h)$; on seeing 1 after h , increment $C(h)$
 - Saturates at boundaries (a saturating counter)
 - Does not decrement below 0 or increment above $2^k - 1$
 - By examining $C(h)$ at any point in time, we can say which outcome had higher likelihood in the last 2^k occurrences of history h
 - Need to shift the origin to mid-point 2^{k-1}
 - If $C(h)$ is below mid-point, the prediction is zero; otherwise 1

30

Direction predictors




31

Direction predictors

- How many different histories? [Size of BHT]
 - No history? [no BHT] Two possible designs:
 - Just a global counter that counts occurrences of 0's and 1's and offers prediction based on that
 - Not very useful
 - One counter per branch (known as bimodal predictor)
 - Somewhat useful: helps identify largely bimodal distributions
 - One global history? [Size of BHT = 1]
 - Captures cross-correlation between branches
 - Since history is a sliding pattern, h will keep on changing
 - One counter for each history pattern? Use a hashmap
 - One local history per branch? Size of BHT?
 - Hash history patterns to counters
 - Loses global correlation


32

Bimodal predictors



33

Bimodal predictors



34

Bimodal predictors

- Simplest of the lot (used in MIPS R10000)
 - Maintains a pattern history table (PHT)
 - MIPS called it a branch history table (BHT)
 - Each entry is a saturating counter
 - Basic idea is to go with the most frequent pattern seen in the past
 - How do you index it?
 - Aliasing?
 - How wide is each counter?
 - Performance of loops (simplest kind of control flow)
 - Alternating branches?
 - Correlating branches?

35

Challenges

- Aliasing in pattern history table
 - Multiple branches mapping to the same counter
- Height of pattern history table is exponential in history length
 - Need long history for better prediction
 - Neural predictors help
- What are the best ways to combine multiple predictors?
 - Each predictor can be an expert for a certain type of branch

36

Pipeline hazard: data

- Pipelining disturbs the sequential thought-process
 - Data dependencies among instructions start to show up
- add \$1, \$2, \$3
 sub \$4, \$1, \$5
 and \$6, \$1, \$7
 or \$8, \$1, \$9
 xor \$10, \$1, \$11
- Result of add is needed by all instructions (RAW hazard: read after write hazard)
- | | | | | | | | | | |
|-----|----|----|----|-----|-----|-----|-----|-----|----|
| add | IF | ID | EX | MEM | WB | | | | |
| sub | | IF | ID | EX | MEM | WB | | | |
| and | | | IF | ID | EX | MEM | WB | | |
| or | | | | IF | ID | EX | MEM | WB | |
| xor | | | | | IF | ID | EX | MEM | WB |
- Red arrows indicate data dependencies from the add instruction to the sub, and, or, and xor instructions.

Pipeline hazard: data

- How to avoid increasing CPI?
 - Stalling is clearly not acceptable
 - Phased register file solves three-cycle apart RAW
 - Can we forward the correct value just in time?
- ```

add IF ID EX MEM WB
sub IF ID EX MEM WB
and IF ID EX MEM WB

```
- Read wrong value in ID/RF, but bypassed value overrides it (need a multiplexor for each ALU input to choose between the RF value and the bypassed value)

38

### Data hazards

- MEM/WB to MEM bypass

```

add r1, r2, r3
lw r4, 0(r1)
sw r4, 20(r1)

```

```

add IF ID EX MEM WB
lw IF ID EX MEM WB
sw IF ID EX MEM WB

```

39

### Pipeline hazard: data

- Can we always avoid stalling?
 

```

lw $1, 0($2)
sub $4, $1, $5
and $6, $1, $7
or $8, $1, $9

```

```

lw IF ID EX MEM WB
sub IF ID EX MEM WB
and IF ID EX MEM WB
or IF ID EX MEM WB

```

  - Need some time travel (backwards)! Not yet feasible!!
  - One option: hardware *pipeline interlock* to stall the *sub* by a cycle
  - Early generations of MIPS (Microprocessor without Interlocked Pipe Stages) had the compiler to fill the *load delay slot* with something independent or a NOP

40

### Hiding data hazard stalls

- In general, cache misses for load instructions can introduce large stalls in the pipeline
- Some of the established ways to hide some of these stalls
  - Execute instructions that do not depend on the load instruction that has missed in the cache
    - Requires ways to reorder the instruction sequence inside the processor
  - Prefetch data into cache before it is needed
    - Need to predict the sequence of data addresses and prefetch from memory hierarchy based on the prediction
  - Predict return value of the load instruction even before it executes

41

### Pipelined instruction execution

- All three types of hazards can cause pipeline stalls and introduce bubbles in the pipeline depending on the pipeline organization
- Overall speedup of pipelining over multi-cycle non-pipelined implementation =  $\text{number of pipe stages} / (1 + \text{average stall cycles per instruction})$

42